

Final Project Submission

Please fill out:

- Student name: Group 5 -
- Student pace: part time
- Scheduled project review date/time:
- Instructor name:
- Blog post URL:

Business Understanding

To develop a predictive model that accurately estimates the sale price of houses in King County, Washington, based on various features such as the number of bedrooms, bathrooms, square footage, location attributes (like waterfront proximity), house condition, grade, year built and renovated. This model aims to assist potential buyers, sellers, and real estate stakeholders in understanding and predicting house prices within the region.

Importing Dependencies

```
In [1]: 1 import itertools
        2 import numpy as np
        3 import pandas as pd
        4 from numbers import Number
        5 from scipy import stats
        6 import matplotlib.pyplot as plt
        7 import seaborn as sns
        8 %matplotlib inline
        9
       10 import pickle
       11
       12 import warnings
       13 warnings.filterwarnings('ignore')
```

Data Understanding

```
In [2]: 1 # Load the dataset
        2 house_data = pd.read_csv("data/kc_house_data.csv")
```

```
In [3]: 1 # structure of dataset
        2 house_data.shape
```

```
Out[3]: (21597, 21)
```

In [4]:

```

1 # Get information about the dataset, including data types and missing v
2 house_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21597 non-null  int64
1   date                 21597 non-null  object
2   price               21597 non-null  float64
3   bedrooms            21597 non-null  int64
4   bathrooms           21597 non-null  float64
5   sqft_living         21597 non-null  int64
6   sqft_lot            21597 non-null  int64
7   floors              21597 non-null  float64
8   waterfront          19221 non-null  object
9   view                21534 non-null  object
10  condition            21597 non-null  object
11  grade               21597 non-null  object
12  sqft_above          21597 non-null  int64
13  sqft_basement       21597 non-null  object
14  yr_built            21597 non-null  int64
15  yr_renovated        17755 non-null  float64
16  zipcode             21597 non-null  int64
17  lat                 21597 non-null  float64
18  long                21597 non-null  float64
19  sqft_living15       21597 non-null  int64
20  sqft_lot15          21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

In [5]:

1

Summary statistics of numerical columns

2

house_data.describe().transpose()

Out[5]:

	count	mean	std	min	25%	50%
id	21597.0	4.580474e+09	2.876736e+09	1.000102e+06	2.123049e+09	3.904930e+09
price	21597.0	5.402966e+05	3.673681e+05	7.800000e+04	3.220000e+05	4.500000e+05
bedrooms	21597.0	3.373200e+00	9.262989e-01	1.000000e+00	3.000000e+00	3.000000e+00
bathrooms	21597.0	2.115826e+00	7.689843e-01	5.000000e-01	1.750000e+00	2.250000e+00
sqft_living	21597.0	2.080322e+03	9.181061e+02	3.700000e+02	1.430000e+03	1.910000e+03
sqft_lot	21597.0	1.509941e+04	4.141264e+04	5.200000e+02	5.040000e+03	7.618000e+03
floors	21597.0	1.494096e+00	5.396828e-01	1.000000e+00	1.000000e+00	1.500000e+00
sqft_above	21597.0	1.788597e+03	8.277598e+02	3.700000e+02	1.190000e+03	1.560000e+03
yr_built	21597.0	1.971000e+03	2.937523e+01	1.900000e+03	1.951000e+03	1.975000e+03
yr_renovated	17755.0	8.363678e+01	3.999464e+02	0.000000e+00	0.000000e+00	0.000000e+00
zipcode	21597.0	9.807795e+04	5.351307e+01	9.800100e+04	9.803300e+04	9.806500e+04
lat	21597.0	4.756009e+01	1.385518e-01	4.715590e+01	4.747110e+01	4.757180e+01
long	21597.0	-1.222140e+02	1.407235e-01	-1.225190e+02	-1.223280e+02	-1.222310e+02
sqft_living15	21597.0	1.986620e+03	6.852305e+02	3.990000e+02	1.490000e+03	1.840000e+03
sqft_lot15	21597.0	1.275828e+04	2.727444e+04	6.510000e+02	5.100000e+03	7.620000e+03

Data Preparation

In [6]:

1

#check for duplicate values

2

house_data.duplicated().sum()

Out[6]: 0

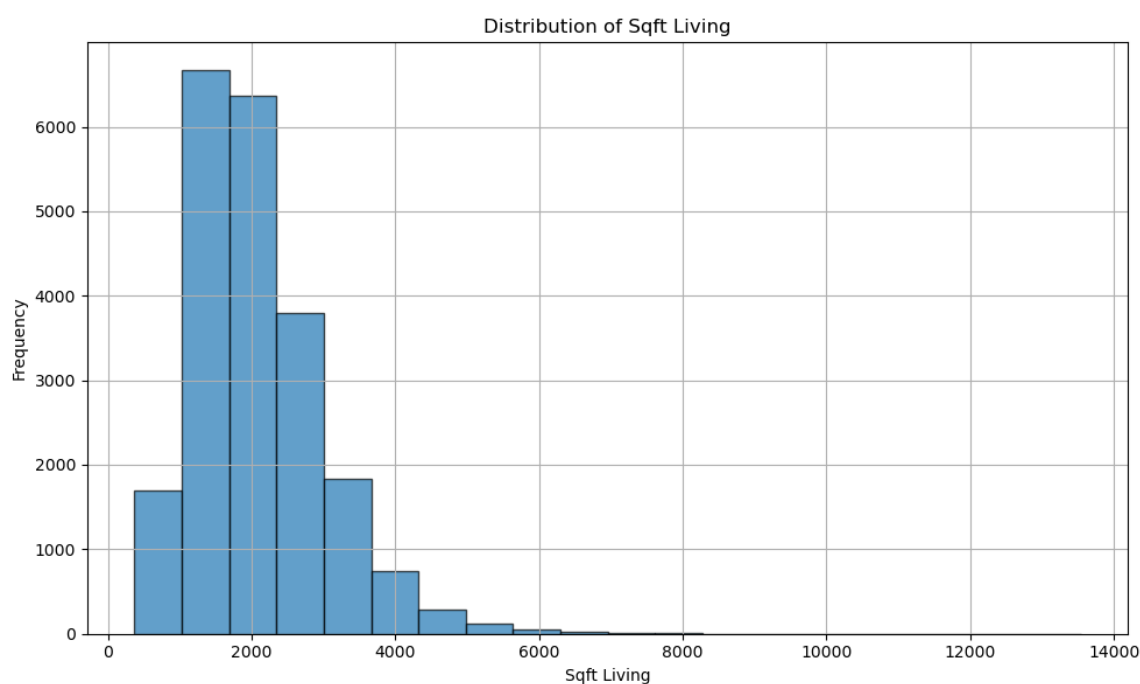
```
In [7]: 1 # Check for missing values in the dataset
        2 house_data.isnull().sum()
```

```
Out[7]: id                0
        date              0
        price             0
        bedrooms          0
        bathrooms         0
        sqft_living       0
        sqft_lot          0
        floors            0
        waterfront       2376
        view              63
        condition         0
        grade             0
        sqft_above        0
        sqft_basement     0
        yr_built          0
        yr_renovated     3842
        zipcode           0
        lat               0
        long              0
        sqft_living15     0
        sqft_lot15        0
        dtype: int64
```

```
In [8]: 1 # fill null values in waterfront column with the mode
        2 house_data['waterfront'].fillna(house_data['waterfront'].mode()[0], inplace=True)
        3
        4 # Convert 'yr_renovated' column from float to integer, handling NaN values
        5 house_data['yr_renovated'] = house_data['yr_renovated'].fillna(0).astype(int)
        6
        7 # Filling missing values for 'yr_renovated' with 0
        8 house_data['yr_renovated'].fillna(house_data['yr_renovated'].median(), inplace=True)
```

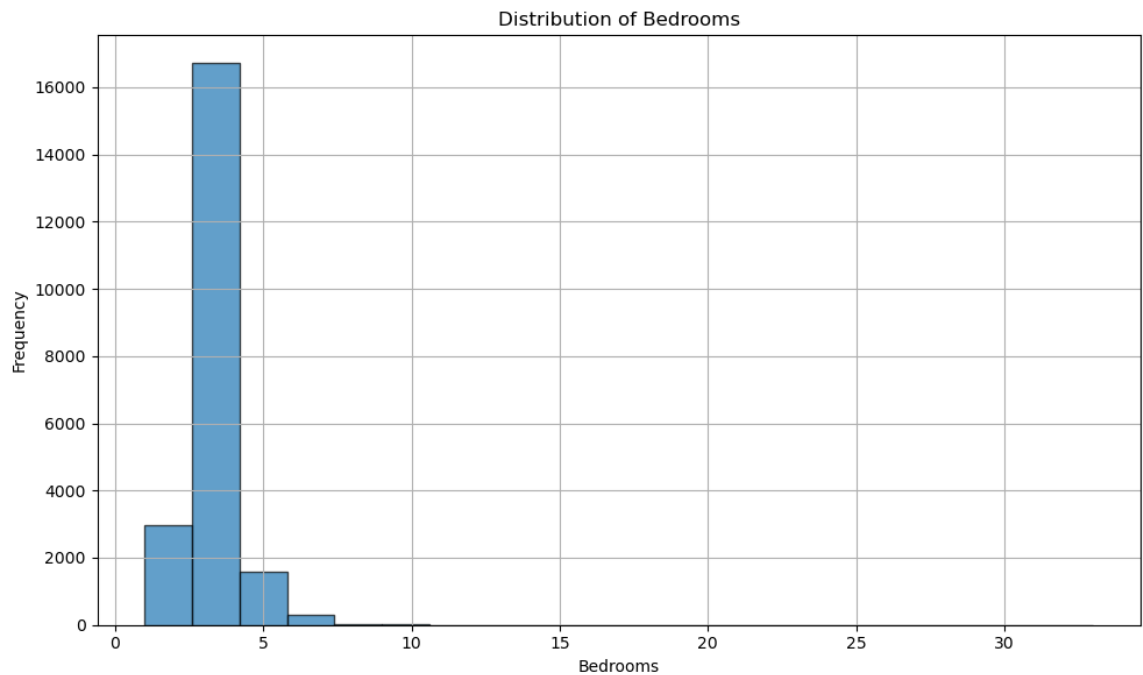
Visualizations

```
In [9]: 1 sqft_living = house_data["sqft_living"]
2
3 # Create the histogram
4 plt.figure(figsize=(10, 6))
5 plt.hist(sqft_living, bins=20, edgecolor="black", alpha=0.7)
6 plt.xlabel("Sqft Living")
7 plt.ylabel("Frequency")
8 plt.title("Distribution of Sqft Living")
9
10 plt.grid(True)
11 plt.tight_layout()
12
13 plt.show()
```



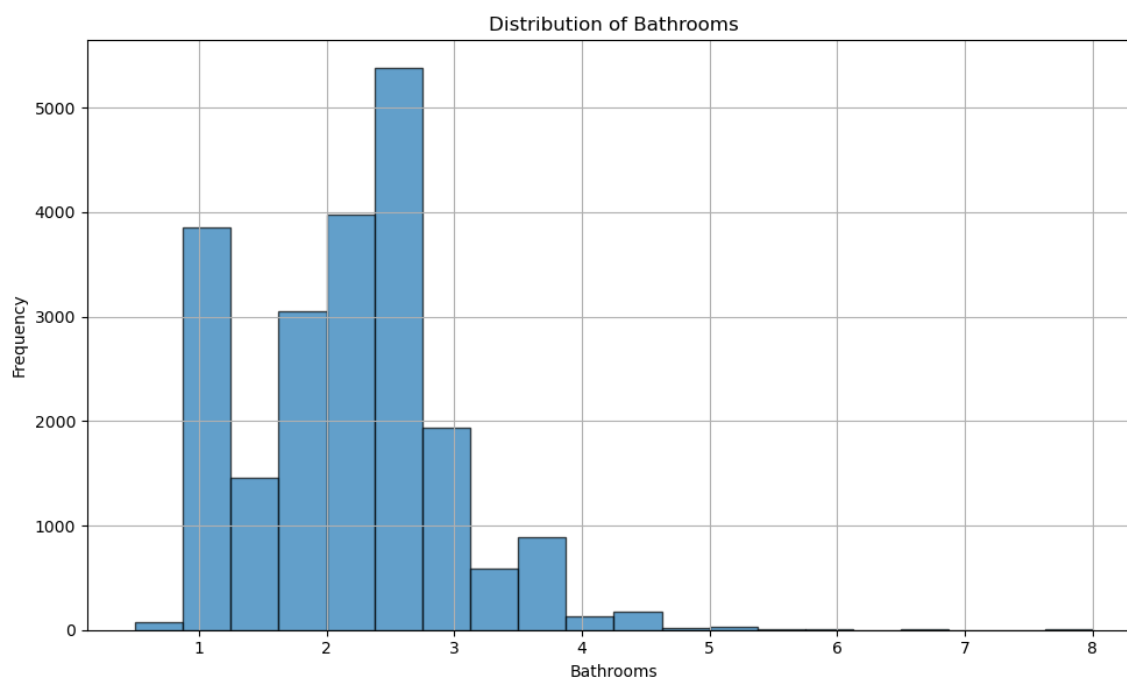
It shows a positive skew, meaning that while most homes fall within a smaller square footage range, there are a few homes with exceptionally large square footage. These larger values, though fewer, extend the tail of the distribution to the right, indicating a positive skew. This suggests that while the majority of homes are smaller, there are a few larger homes that are exceptions.

```
In [10]: 1 bedrooms = house_data['bedrooms']
2
3 # Create the histogram
4 plt.figure(figsize=(10, 6))
5 plt.hist(bedrooms, bins=20, edgecolor="black", alpha=0.7)
6 plt.xlabel("Bedrooms")
7 plt.ylabel("Frequency")
8 plt.title("Distribution of Bedrooms")
9
10 plt.grid(True)
11 plt.tight_layout()
12
13 plt.show()
```



There is a significant peak at 2 to 3 bedrooms, indicating that most homes have this number of bedrooms.


```
In [11]: 1 #checking the distribution of number of bathrooms
2
3 bathrooms = house_data['bathrooms']
4
5 # Create the histogram
6 plt.figure(figsize=(10, 6))
7 plt.hist(bathrooms, bins=20, edgecolor="black", alpha=0.7)
8 plt.xlabel("Bathrooms")
9 plt.ylabel("Frequency")
10 plt.title("Distribution of Bathrooms")
11
12 plt.grid(True)
13 plt.tight_layout()
14
15 plt.show()
```



Most Common: Houses with 3 bathrooms have the highest frequency.

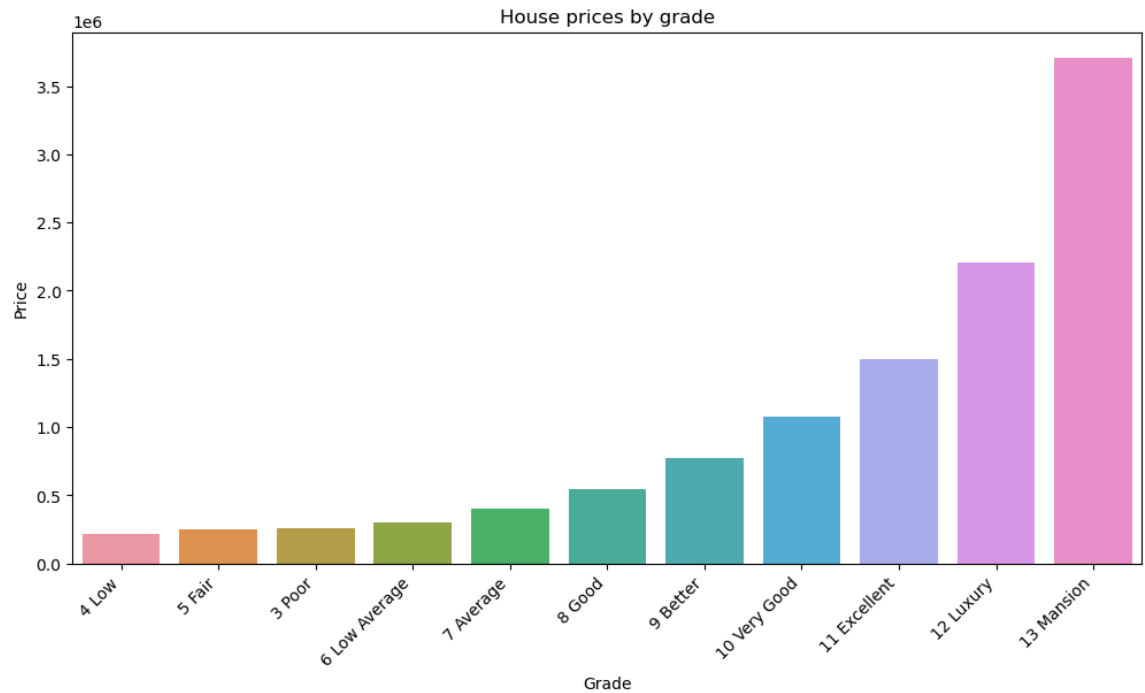
Lower Frequencies: Counts beyond 4 bathrooms have much lower frequencies, suggesting they are less common.

```
In [12]: 1 # Pie chart to show number of houses by condition
2 house_counts = house_data['condition'].value_counts()
3 plt.figure(figsize=(6, 6))
4 explode = (0, 0, 0, 0, 0.3)
5 house_counts.plot(kind='pie', autopct="%1.1f%", explode=explode, title
6 plt.axis('equal')
7 plt.show()
```



Most of the houses are in average to good condition.


```
In [13]: 1 # Distribution of houses by grade
2 ave_prices = house_data[['grade', 'price']].groupby('grade').mean().sort
3 plt.figure(figsize=(12, 6))
4 sns.barplot(x='grade', y='price', data=ave_prices, label='Price')
5 plt.xticks(rotation=45, ha='right')
6 plt.xlabel("Grade")
7 plt.ylabel("Price")
8 plt.title('House prices by grade')
9 plt.show()
10
```



Houses under average

Feature Engineering

```
In [14]: 1 # Explore specific columns related to the features mentioned in the bus
2 relevant_columns = ['date', 'price', 'bedrooms', 'bathrooms', 'floors',
3                   'waterfront', 'condition', 'grade', 'yr_built', 'yr
4 relevant_data = house_data[relevant_columns]
5 relevant_data
```

```
Out[14]:
```

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	co
0	10/13/2014	221900.0	3	1.00	1.0	1180	5650	NO	
1	12/9/2014	538000.0	3	2.25	2.0	2570	7242	NO	
2	2/25/2015	180000.0	2	1.00	1.0	770	10000	NO	
3	12/9/2014	604000.0	4	3.00	1.0	1960	5000	NO	
4	2/18/2015	510000.0	3	2.00	1.0	1680	8080	NO	
...
21592	5/21/2014	360000.0	3	2.50	3.0	1530	1131	NO	
21593	2/23/2015	400000.0	4	2.50	2.0	2310	5813	NO	
21594	6/23/2014	402101.0	2	0.75	2.0	1020	1350	NO	
21595	1/16/2015	400000.0	3	2.50	2.0	1600	2388	NO	
21596	10/15/2014	325000.0	2	0.75	2.0	1020	1076	NO	

21597 rows × 12 columns



```
In [15]: 1 # identify unique values in the columns
2 print(house_data["waterfront"].unique())
3 print()
4 print(house_data["condition"].unique())
5 print()
6 print(house_data["grade"].unique())
```

```
['NO' 'YES']
```

```
['Average' 'Very Good' 'Good' 'Poor' 'Fair']
```

```
['7 Average' '6 Low Average' '8 Good' '11 Excellent' '9 Better' '5 Fair'
 '10 Very Good' '12 Luxury' '4 Low' '3 Poor' '13 Mansion']
```

In [16]:

1

Create a new column indicating if renovation happened (1) or not (0)

2

relevant_data['renovated'] = (relevant_data['yr_renovated'] != 0).astype

3

relevant_data

Out[16]:

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	co
0	10/13/2014	221900.0	3	1.00	1.0	1180	5650	NO	
1	12/9/2014	538000.0	3	2.25	2.0	2570	7242	NO	
2	2/25/2015	180000.0	2	1.00	1.0	770	10000	NO	
3	12/9/2014	604000.0	4	3.00	1.0	1960	5000	NO	
4	2/18/2015	510000.0	3	2.00	1.0	1680	8080	NO	
...	
21592	5/21/2014	360000.0	3	2.50	3.0	1530	1131	NO	
21593	2/23/2015	400000.0	4	2.50	2.0	2310	5813	NO	
21594	6/23/2014	402101.0	2	0.75	2.0	1020	1350	NO	
21595	1/16/2015	400000.0	3	2.50	2.0	1600	2388	NO	
21596	10/15/2014	325000.0	2	0.75	2.0	1020	1076	NO	

21597 rows × 13 columns

```

In [17]: 1 # Convert 'date' column to datetime
2 relevant_data['date'] = pd.to_datetime(relevant_data['date'])
3
4 # Define a function to get the season based on the month
5 def get_season(month):
6     if month in [3, 4, 5]: # Spring: March, April, May
7         return 'Spring'
8     elif month in [6, 7, 8]: # Summer: June, July, August
9         return 'Summer'
10    elif month in [9, 10, 11]: # Fall: September, October, November
11        return 'Fall'
12    else: # Winter: December, January, February
13        return 'Winter'
14
15 # Create 'season' column based on 'date'
16 relevant_data['season'] = relevant_data['date'].dt.month.apply(get_season)
17 relevant_data.head(15)

```

```

Out[17]:

```

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	condition
0	2014-10-13	221900.0	3	1.00	1.0	1180	5650	NO	Average
1	2014-12-09	538000.0	3	2.25	2.0	2570	7242	NO	Average
2	2015-02-25	180000.0	2	1.00	1.0	770	10000	NO	Average
3	2014-12-09	604000.0	4	3.00	1.0	1960	5000	NO	Very Good
4	2015-02-18	510000.0	3	2.00	1.0	1680	8080	NO	Average
5	2014-05-12	1230000.0	4	4.50	1.0	5420	101930	NO	Average
6	2014-06-27	257500.0	3	2.25	2.0	1715	6819	NO	Average
7	2015-01-15	291850.0	3	1.50	1.0	1060	9711	NO	Average
8	2015-04-15	229500.0	3	1.00	1.0	1780	7470	NO	Average
9	2015-03-12	323000.0	3	2.50	2.0	1890	6560	NO	Average
10	2015-04-03	662500.0	3	2.50	1.0	3560	9796	NO	Average
11	2014-05-27	468000.0	2	1.00	1.0	1160	6000	NO	Good
12	2014-05-28	310000.0	3	1.00	1.5	1430	19901	NO	Good
13	2014-10-07	400000.0	3	1.75	1.0	1370	9680	NO	Good
14	2015-03-12	530000.0	5	2.00	1.5	1810	4850	NO	Average

In [18]:

1

Replace 'No' with 0 and 'Yes' with 1 in the 'waterfront' column

2

relevant_data['waterfront'] = relevant_data['waterfront'].replace({'NO'

3

relevant_data

Out[18]:

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	conditi
0	2014-10-13	221900.0	3	1.00	1.0	1180	5650	0	Avera
1	2014-12-09	538000.0	3	2.25	2.0	2570	7242	0	Avera
2	2015-02-25	180000.0	2	1.00	1.0	770	10000	0	Avera
3	2014-12-09	604000.0	4	3.00	1.0	1960	5000	0	Vl Gc
4	2015-02-18	510000.0	3	2.00	1.0	1680	8080	0	Avera
...	
21592	2014-05-21	360000.0	3	2.50	3.0	1530	1131	0	Avera
21593	2015-02-23	400000.0	4	2.50	2.0	2310	5813	0	Avera
21594	2014-06-23	402101.0	2	0.75	2.0	1020	1350	0	Avera
21595	2015-01-16	400000.0	3	2.50	2.0	1600	2388	0	Avera
21596	2014-10-15	325000.0	2	0.75	2.0	1020	1076	0	Avera

21597 rows × 14 columns

In [19]:

```
1 # Extract grade number and condition using regex
2 # Add grade number (grade_no) column
3 import re
4
5 relevant_data[['grade_no', 'condition']] = relevant_data['grade'].str.e
6 relevant_data
```

Out[19]:

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	conditi
0	2014-10-13	221900.0	3	1.00	1.0	1180	5650	0	Avera
1	2014-12-09	538000.0	3	2.25	2.0	2570	7242	0	Avera
2	2015-02-25	180000.0	2	1.00	1.0	770	10000	0	L Avera
3	2014-12-09	604000.0	4	3.00	1.0	1960	5000	0	Avera
4	2015-02-18	510000.0	3	2.00	1.0	1680	8080	0	Gc
...	
21592	2014-05-21	360000.0	3	2.50	3.0	1530	1131	0	Gc
21593	2015-02-23	400000.0	4	2.50	2.0	2310	5813	0	Gc
21594	2014-06-23	402101.0	2	0.75	2.0	1020	1350	0	Avera
21595	2015-01-16	400000.0	3	2.50	2.0	1600	2388	0	Gc
21596	2014-10-15	325000.0	2	0.75	2.0	1020	1076	0	Avera

21597 rows × 15 columns



In [20]:

```
1 # convert grade_no values to int
2 relevant_data['grade_no'] = relevant_data['grade_no'].astype('int64')
```

In [21]:

```
1 print(relevant_data.columns)
```

Index(['date', 'price', 'bedrooms', 'bathrooms', 'floors', 'sqft_living',
 'sqft_lot', 'waterfront', 'condition', 'grade', 'yr_built',
 'yr_renovated', 'renovated', 'season', 'grade_no'],
 dtype='object')

In [22]:

1

2

3

4

5

Create a new feature for the age of the house

current_year = pd.Timestamp.now().year # Get the current year

relevant_data['house_age'] = current_year - relevant_data['yr_built']

relevant_data.head(15)

Out[22]:

	date	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	condition
0	2014-10-13	221900.0	3	1.00	1.0	1180	5650	0	Average
1	2014-12-09	538000.0	3	2.25	2.0	2570	7242	0	Average
2	2015-02-25	180000.0	2	1.00	1.0	770	10000	0	Low Average
3	2014-12-09	604000.0	4	3.00	1.0	1960	5000	0	Average
4	2015-02-18	510000.0	3	2.00	1.0	1680	8080	0	Good
5	2014-05-12	1230000.0	4	4.50	1.0	5420	101930	0	Excellent
6	2014-06-27	257500.0	3	2.25	2.0	1715	6819	0	Average
7	2015-01-15	291850.0	3	1.50	1.0	1060	9711	0	Average
8	2015-04-15	229500.0	3	1.00	1.0	1780	7470	0	Average
9	2015-03-12	323000.0	3	2.50	2.0	1890	6560	0	Average
10	2015-04-03	662500.0	3	2.50	1.0	3560	9796	0	Good
11	2014-05-27	468000.0	2	1.00	1.0	1160	6000	0	Average
12	2014-05-28	310000.0	3	1.00	1.5	1430	19901	0	Average
13	2014-10-07	400000.0	3	1.75	1.0	1370	9680	0	Average
14	2015-03-12	530000.0	5	2.00	1.5	1810	4850	0	Average

In [23]:

1

Drop irrelevant columns and assign the result back to engineered_data

2

relevant_data = relevant_data.drop(['date', 'grade', 'yr_renovated', 'y

3

relevant_data

Out[23]:

	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	renovated	s
0	221900.0	3	1.00	1.0	1180	5650	0	0	
1	538000.0	3	2.25	2.0	2570	7242	0	1	
2	180000.0	2	1.00	1.0	770	10000	0	0	
3	604000.0	4	3.00	1.0	1960	5000	0	0	
4	510000.0	3	2.00	1.0	1680	8080	0	0	
...	
21592	360000.0	3	2.50	3.0	1530	1131	0	0	
21593	400000.0	4	2.50	2.0	2310	5813	0	0	
21594	402101.0	2	0.75	2.0	1020	1350	0	0	St
21595	400000.0	3	2.50	2.0	1600	2388	0	0	
21596	325000.0	2	0.75	2.0	1020	1076	0	0	

21597 rows × 11 columns

Data Visualizations

In [24]:

1

Pairplot

2

sns.pairplot(relevant_data)

3

plt.show()

In [25]:


1

relevant_data.price.describe()

Out[25]:

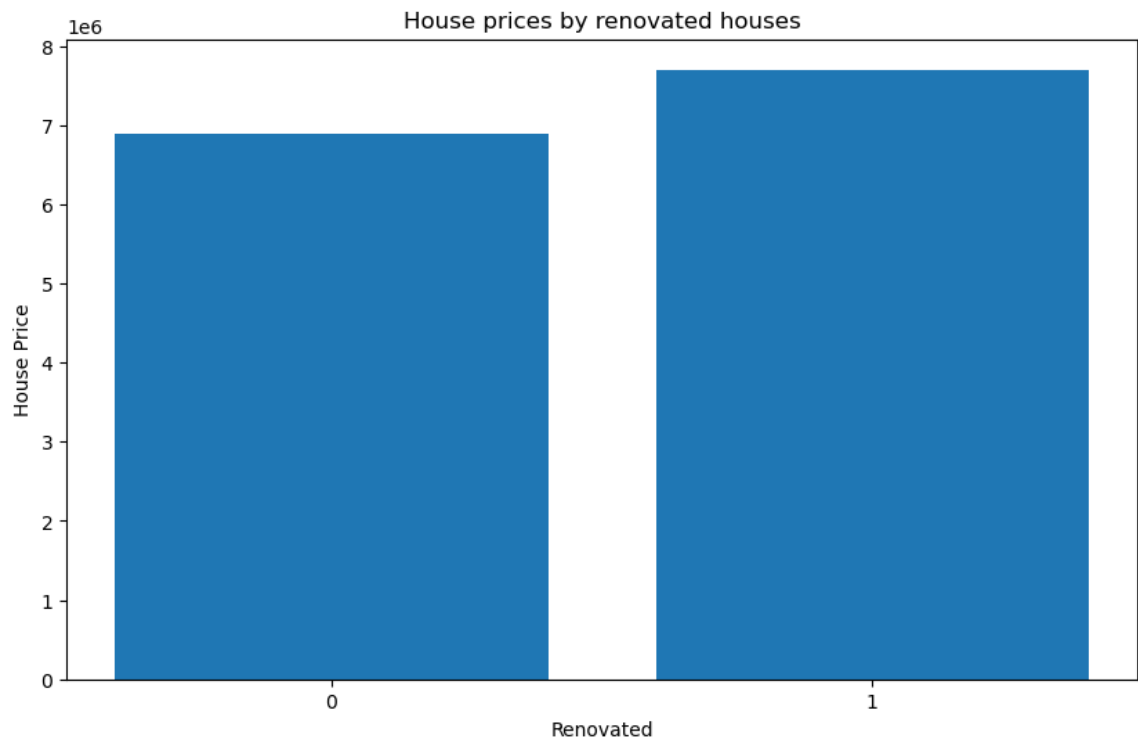
count	2.159700e+04
mean	5.402966e+05
std	3.673681e+05
min	7.800000e+04
25%	3.220000e+05
50%	4.500000e+05
75%	6.450000e+05
max	7.700000e+06
Name: price, dtype: float64	


```
In [26]: 1 # check the average house prices by season
2 average_prices = relevant_data.groupby('season')['price'].mean().reset_
3
4 # Bar plot using Seaborn with average prices
5 plt.figure(figsize=(8, 4))
6 sns.barplot(x='season', y='price', data=average_prices)
7 plt.title('Average House Prices by Seasons')
8 plt.xlabel('Seasons')
9 plt.ylabel('Average Prices')
10 plt.savefig('house_grade.png', transparent=True)
11 plt.show()
```




Average house prices are highest during spring and slightly lower during winter.

```
In [27]: 1 # checking the distribution of house prices by renovated houses
2 Renovated = relevant_data['renovated']
3 Price = relevant_data['price']
4 plt.figure(figsize=(10,6))
5 plt.bar(Renovated, Price)
6 plt.xlabel('Renovated')
7 plt.ylabel('House Price')
8 plt.title('House prices by renovated houses')
9 plt.xticks(ticks=[0, 1])
10 plt.show()
```



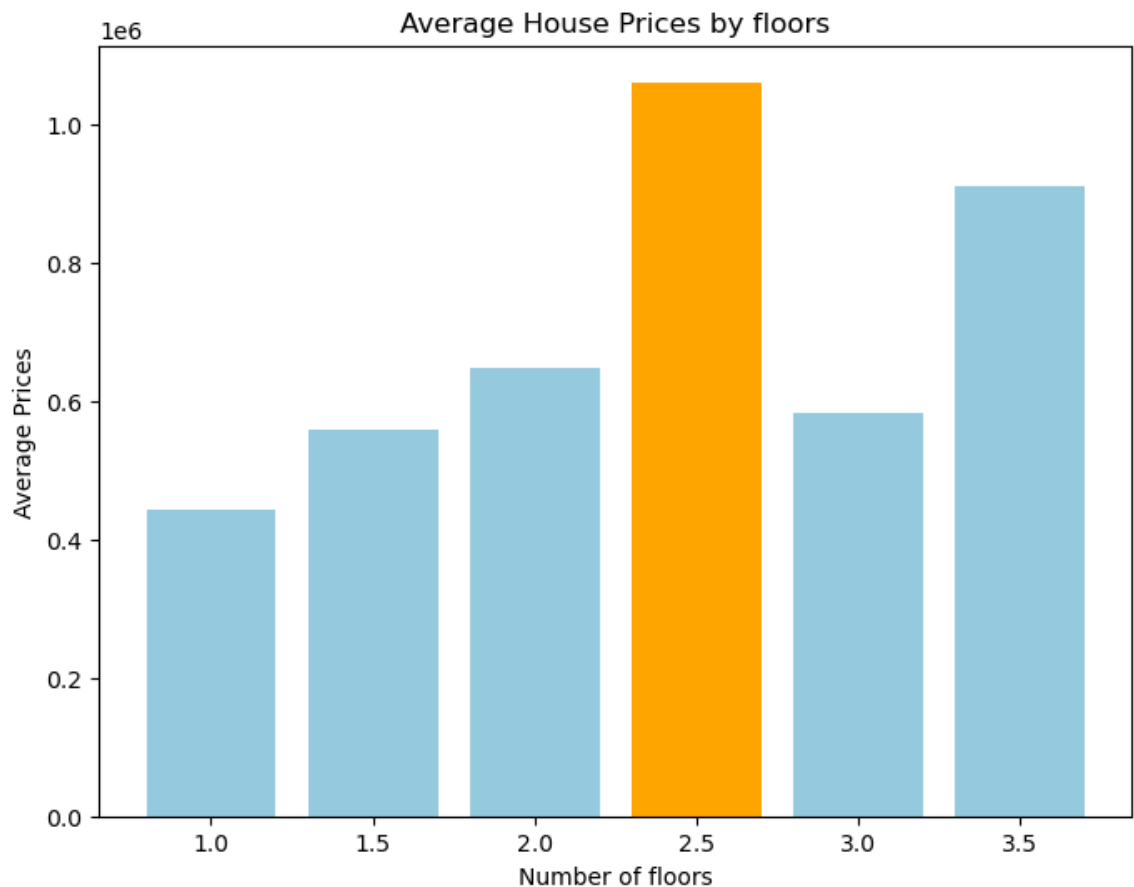
Renovated houses have a higher average price indicating renovations had a small impact on the value of the house.

```
In [28]: 1 # Distribution of house price by age
2 House_age = relevant_data['house_age']
3 Price = relevant_data['price']
4 plt.figure(figsize=(12,6))
5 plt.bar(House_age, Price)
6 plt.xlabel('House Age in Years')
7 plt.ylabel('House Price')
8 plt.title('House prices by house age')
9
10 plt.show()
```



The houses are priced at 2 million dollars and above where the age is below 50 years. On the contrary majority of the houses over 50 years are priced below 2 million dollars.


```
In [29]: 1 # Distribution of houses by floors
2 # Get the number of houses with each number of floors
3 average_prices = relevant_data.groupby('floors')['price'].mean().reset_
4
5 # Find the index of the maximum value
6 max_index = average_prices['price'].idxmax()
7
8 # Bar plot using Seaborn with average prices
9 plt.figure(figsize=(8, 6))
10 sns.barplot(x='floors', y='price', data=average_prices, color='skyblue')
11
12 # Highlight the bar with the maximum value
13 plt.bar(max_index, average_prices.loc[max_index, 'price'], color='orang
14
15 plt.title('Average House Prices by floors')
16 plt.xlabel('Number of floors')
17 plt.ylabel('Average Prices')
18
19 plt.show()
```




Houses with 2.5 floors are the most priced at slightly above 1 million dollars and the least priced houses have 1 floor.

Checking for outliers

```
In [30]: 1 # seasons and price  
2 sns.boxplot(x='season', y='price', data=relevant_data)  
3 plt.show()
```




```
In [31]: 1 # bedrooms and price
2 sns.boxplot(x=relevant_data['bedrooms'], y=relevant_data['price'])
3 plt.show()
```



Removing outliers using z-score

```
In [32]: 1 plt.figure(figsize=(6, 2))
2 sns.boxplot(x='bedrooms', data=relevant_data)
3 plt.title('Boxplot of bedrooms')
4 plt.show()
```




```
In [33]: 1 #removing outlier with '33' bedrooms
2 relevant_data=relevant_data[relevant_data['bedrooms']!=33]
```


```
In [34]: 1 # Calculate Z-score for 'price'
2 z_scores = stats.zscore(relevant_data['price'])
3 threshold = 3 # Define a threshold for extreme values (e.g., 3 standar
4
5 # Detect outliers
6 outliers = relevant_data['price'][np.abs(z_scores) > threshold]
7
8 # Remove outliers
9 relevant_data1 = relevant_data[np.abs(z_scores) <= threshold]
```

```
In [35]: 1 # Pairplot
2 # sns.pairplot(relevant_data)
3 # plt.show()
```


```
In [36]: 1 # seasons and price
2 sns.boxplot(x='season', y='price', data=relevant_data)
3 plt.show()
```



```
In [37]: 1 sns.boxplot(x='bedrooms', y='price', data=relevant_data)
          2 plt.show()
```




```
In [38]: 1 #create a scatter plot for the numeric columns to observe their relation
2 numeric_columns = relevant_data.select_dtypes(include=['number'])
3
4 # Iterate through each numeric column and create scatter plots
5 for column in numeric_columns.columns:
6     plt.scatter(relevant_data[column], relevant_data['price'], alpha=0.
7     plt.title(f'Scatter Plot of {column} against Price')
8     plt.xlabel(column)
9     plt.ylabel('Price')
10    plt.grid(True)
11    plt.show()
```



Correlation

```
In [39]: 1 relevant_data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21596 entries, 0 to 21596
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   price           21596 non-null  float64
 1   bedrooms        21596 non-null  int64  
 2   bathrooms       21596 non-null  float64
 3   floors          21596 non-null  float64
 4   sqft_living     21596 non-null  int64  
 5   sqft_lot        21596 non-null  int64  
 6   waterfront      21596 non-null  int64  
 7   renovated       21596 non-null  int32  
 8   season          21596 non-null  object  
 9   grade_no        21596 non-null  int64  
10  house_age       21596 non-null  int64  
dtypes: float64(3), int32(1), int64(6), object(1)
memory usage: 1.9+ MB
```

```
In [40]: 1 numeric_df = relevant_data.select_dtypes(include=['float64', 'int64'])
2         corr_price = numeric_df.corr()['price'].sort_values()
3         corr_price
```

```
Out[40]: house_age      -0.053965
sqft_lot      0.089879
floors        0.256820
waterfront    0.264308
bedrooms      0.315961
bathrooms     0.525915
grade_no      0.667964
sqft_living   0.701929
price         1.000000
Name: price, dtype: float64
```

```
In [41]: 1 relevant_data.columns
```


```
Out[41]: Index(['price', 'bedrooms', 'bathrooms', 'floors', 'sqft_living', 'sqft_lo
t',
               'waterfront', 'renovated', 'season', 'grade_no', 'house_age'],
              dtype='object')
```

```
In [42]: 1 # Correlation table
2         cols = ['price', 'bedrooms', 'bathrooms', 'floors', 'sqft_living', 'sqf
3             'waterfront', 'renovated', 'grade_no', 'house_age']
4         corr = relevant_data[cols].corr()
5         corr
```

```
Out[42]:
```

	price	bedrooms	bathrooms	floors	sqft_living	sqft_lot	waterfront	ren
price	1.000000	0.315961	0.525915	0.256820	0.701929	0.089879	0.264308	0
bedrooms	0.315961	1.000000	0.527870	0.183707	0.593178	0.033602	-0.002054	0.
bathrooms	0.525915	0.527870	1.000000	0.502574	0.755755	0.088368	0.063628	0.
floors	0.256820	0.183707	0.502574	1.000000	0.353941	-0.004824	0.020794	0.
sqft_living	0.701929	0.593178	0.755755	0.353941	1.000000	0.173449	0.104635	0.
sqft_lot	0.089879	0.033602	0.088368	-0.004824	0.173449	1.000000	0.021458	0.
waterfront	0.264308	-0.002054	0.063628	0.020794	0.104635	0.021458	1.000000	0.
renovated	0.117546	0.018354	0.046738	0.003705	0.050825	0.005089	0.074267	1.
grade_no	0.667964	0.366174	0.665834	0.458783	0.762776	0.114726	0.082817	0.
house_age	-0.053965	-0.160736	-0.507166	-0.489175	-0.318140	-0.052939	0.024491	0.

```
In [43]: 1 # heatmap with seaborn
2 sns.set(font_scale=0.8)
3 sns.heatmap(corr, annot=True, cmap='Reds')
4 plt.title("Correlation Heatmap")
5 plt.show()
```



- The strongest positive correlation (0.7) is between price and square footage of living area, indicating that larger homes tend to have higher prices.
- The second strongest positive correlation (0.67) is between price and the grade_no, suggesting that the higher the grade the higher the price.
- The third strongest positive correlation (0.32) is between price and number of bedrooms, implying that more bedrooms in a home also contribute to its price, but to a lesser extent than bathrooms.
- The weakest positive correlation (0.05) is between price and year built, meaning that newer homes have slightly higher prices than older homes, but the effect is negligible.
- Price and Year Renovated is only (0.12). Reasons could be:
 - Renovations might cater to personal preferences rather than broad market appeal.
 - Recent renovations might indicate potential underlying issues with the property.

Modelling

```
In [44]: 1 import statsmodels.formula.api as sm
```

In [45]:

```
1  # To build a regression model we start with the baseline model.
2  # This is a simple linear regression
3  # We start by assigning the dependent variable and the independent vari
4
5  y = relevant_data['price']
6  x_baseline = relevant_data[['sqft_living']]
```

```
In [46]: 1 # Build the simple linear regression model and display the results of the
2 import statsmodels.api as sm
3
4 baseline_model = sm.OLS(y, sm.add_constant(x_baseline))
5 baseline_results = baseline_model.fit()
6 print(baseline_results.summary())
```

OLS Regression Results

```
=====
===
Dep. Variable:          price    R-squared:
0.493
Model:                OLS    Adj. R-squared:
0.493
Method:             Least Squares    F-statistic:          2.097
e+04
Date:                Wed, 03 Jan 2024    Prob (F-statistic):
0.00
Time:                21:44:56    Log-Likelihood:          -3.0005
e+05
No. Observations:    21596    AIC:          6.001
e+05
Df Residuals:        21594    BIC:          6.001
e+05
Df Model:              1
Covariance Type:      nonrobust
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const      -4.401e+04    4410.123     -9.980     0.000    -5.27e+04    -3.5
4e+04
sqft_living    280.8688      1.939     144.820     0.000     277.067     28
4.670
=====
===
Omnibus:          14801.492    Durbin-Watson:
1.982
Prob(Omnibus):      0.000    Jarque-Bera (JB):          54264
2.481
Skew:              2.820    Prob(JB):
0.00
Kurtosis:          26.901    Cond. No.          5.63
e+03
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [47]: 1 # Plot the scatter diagram and the best_fit regression line.
2
3 fig, ax = plt.subplots()
4 relevant_data.plot.scatter(x = 'sqft_living', y = 'price', label = 'Data points')
5 sm.graphics.abline_plot(model_results = baseline_results, label = 'regression line')
6 ax.legend();
```



Intepretation of the Baseline model

- The R squared of the model is 0.49 interpreted as the propotional extent to which independent variable, Squarefoot of living area, explains the change in price.
- The probability of the F Statistic is well below the alpha value of 0.05 which shows that model is stastically significant.
- The coefficients of the model (const (which is the y intercept) and the sqft_living) are both statistically significant with a p value well below 0.05.
- For each increase in sqft_living we see a corresponding increase of 281 in price.

Building a Multiple regression Model

```
In [48]: 1 relevant_data.columns
```

```
Out[48]: Index(['price', 'bedrooms', 'bathrooms', 'floors', 'sqft_living', 'sqft_lot',
               'waterfront', 'renovated', 'season', 'grade_no', 'house_age'],
              dtype='object')
```

In [49]:

```
1 y = relevant_data['price']  
2 x_baseline = relevant_data[['grade_no', 'sqft_living', 'bathrooms', 'renov
```

In [50]:

```
1  # Build the simple linear regression model and display the results of t  
2  import statsmodels.api as sm  
3  
4  
5  baseline_model = sm.OLS(y, sm.add_constant(x_baseline))  
6  baseline_results = baseline_model.fit()  
7  print(baseline_results.summary())
```



```

=====
OLS Regression Results
=====
=====
Dep. Variable:                price    R-squared:
0.607
Model:                        OLS      Adj. R-squared:
0.607
Method:                      Least Squares    F-statistic:                6
672.
Date:                        Wed, 03 Jan 2024    Prob (F-statistic):
0.00
Time:                        21:44:58    Log-Likelihood:                -2.9729
e+05
No. Observations:            21596    AIC:                5.946
e+05
Df Residuals:                21590    BIC:                5.946
e+05
Df Model:                    5
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const      -1.166e+06    1.55e+04    -75.151    0.000    -1.2e+06    -1.1
4e+06
grade_no    1.417e+05    2200.325     64.409    0.000    1.37e+05    1.4
6e+05
sqft_living  155.1451         3.176     48.843    0.000    148.919     16
1.371
bathrooms   4.247e+04    3484.316     12.188    0.000    3.56e+04     4.9
3e+04
renovated    4.715e+04    8921.922      5.285    0.000    2.97e+04     6.4
6e+04
house_age    3907.1281      67.185     58.155    0.000    3775.440    403
8.816
=====
=====
Omnibus:                17739.221    Durbin-Watson:
1.977
Prob(Omnibus):          0.000    Jarque-Bera (JB):                133922
6.121
Skew:                   3.466    Prob(JB):
0.00
Kurtosis:               40.951    Cond. No.                2.28
e+04
=====
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is cor
rectly specified.
[2] The condition number is large, 2.28e+04. This might indicate that ther
e are
strong multicollinearity or other numerical problems.
```

```
In [51]: 1 y = relevant_data['price']  
2 x_baseline = relevant_data[['bedrooms', 'bathrooms', 'floors', 'sqft_li  
3 'waterfront', 'renovated', 'grade_no', 'house_age']]
```

```
In [52]: 1 # Build the simple linear regression model and display the results of t  
2 baseline_model = sm.OLS(y, sm.add_constant(x_baseline))  
3 baseline_results = baseline_model.fit()  
4 print(baseline_results.summary())
```

OLS Regression Results

```
=====
====
Dep. Variable:                price    R-squared:
0.646
Model:                        OLS      Adj. R-squared:
0.645
Method:                      Least Squares    F-statistic:                4
368.
Date:                        Wed, 03 Jan 2024    Prob (F-statistic):
0.00
Time:                        21:44:58    Log-Likelihood:                -2.9618
e+05
No. Observations:            21596    AIC:                5.924
e+05
Df Residuals:                21586    BIC:                5.925
e+05
Df Model:                    9
Covariance Type:            nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]

const	-1.017e+06	1.61e+04	-63.270	0.000	-1.05e+06	-9.8
5e+05						
bedrooms	-4.573e+04	2134.899	-21.418	0.000	-4.99e+04	-4.1
5e+04						
bathrooms	5.282e+04	3477.102	15.192	0.000	4.6e+04	5.9
6e+04						
floors	1.692e+04	3436.267	4.923	0.000	1.02e+04	2.3
7e+04						
sqft_living	179.8298	3.324	54.106	0.000	173.315	18
6.344						
sqft_lot	-0.2495	0.037	-6.781	0.000	-0.322	-
0.177						
waterfront	7.512e+05	1.84e+04	40.838	0.000	7.15e+05	7.8
7e+05						
renovated	1.759e+04	8509.590	2.067	0.039	907.193	3.4
3e+04						
grade_no	1.291e+05	2159.160	59.791	0.000	1.25e+05	1.3
3e+05						
house_age	3957.4099	66.448	59.557	0.000	3827.168	408
7.652						

```
=====
====
Omnibus:                    15650.960    Durbin-Watson:
1.976
Prob(Omnibus):              0.000    Jarque-Bera (JB):                96655
6.213
Skew:                      2.888    Prob(JB):
0.00
Kurtosis:                  35.261    Cond. No.                5.45
e+05
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, $5.45e+05$. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation of results

- R-squared: 0.646 Indicates that the model explains 64.6 % of the variance in the dependent variable (price).
- Adj. R-squared: 0.645 Similar to R-squared but adjusts for the number of predictors in the model.
- Prob (F-statistic): 0.00 The probability associated with the F-statistic, suggesting the overall significance of the model.
- Coefficients: Each predictor variable's coefficient represents the change in the dependent variable (price) per unit change in that predictor, holding other predictors constant.
 - Intercept (const): $-1.017e+06$
 - bedrooms: $-4.573e+04$
 - bathrooms: $5.282e+04$
 - floors: $1.692e+04$
 - sqft_living: 179.8298
 - sqft_lot: -0.2495
 - waterfront: $7.512e+05$
 - renovated: $1.759e+04$
 - grade_no: $1.291e+05$
 - house_age: 3957.4099
- Statistical Significance (P-values): $P > |t|$ values for each coefficient indicate the probability of observing the data if the null hypothesis (that the coefficient is zero) is true.
- Variables with a P-value less than a significance level (commonly 0.05) are typically considered statistically significant. For instance, bathrooms, floors, sqft_living, sqft_lot, waterfront, renovated, grade_no, house_age, seem significant since their P-values are 0.000.

Recommendations

Our potential homeowners are advised to focus on the variables on the final model when looking for competitively priced homes. These are majorly square feet of the living space, bathrooms and bedrooms. Where one is looking to buy a home at favorable prices a potential homeowner will need to compromise on one or two items. e.g. waterfront homes or square feet of living area. Our stakeholders are advised to purchase homes in the spring or summer in order to get a good variety of homes to pick from. Further Data Collection: Explore additional data beyond the current predictors to refine the model. Factors like proximity to amenities, neighborhood trends, or specific architectural styles could further enhance price prediction accuracy.